


A Guide to Julia's Dependencies

aka why does Julia take so long to compile from source?

Tony Kelman

(@tkelman on Github)



Bay Area Julia Users

December 4, 2014

Who's this guy?

- Grad student at Berkeley in Mechanical Engineering
- Julia user and contributor since February 2014



(look up "Rejected" on Youtube)

- On Github:
- Because I was curious, which files have I committed to most?

```
$ git log --author=kelman --pretty=oneline --name-only | sort | uniq -c | sort -r -n | head -n 15
29 deps/Makefile
17 Make.inc
12 Makefile
8 base/interactiveutil.jl
6 src/flisp/Makefile
6 README.md
5 test/sparse.jl
5 test/file.jl
5 src/sys.c
5 src/support/Makefile
5 src/debuginfo.cpp
5 LICENSE.md
5 contrib/windows/msys_build.sh
5 appveyor.yml
5 .travis.yml
```

Who's this guy?

- Grad student at Berkeley in Mechanical Engineering
- Julia user and contributor since February 2014



(look up "Rejected" on Youtube)

- On Github:
- Because I was curious, which files have I committed to most?

```
$ git log --author=kelman --pretty=oneline --name-only | sort | uniq -c | sort -r -n | head -n 15
 29 deps/Makefile
 17 Make.inc
 12 Makefile
   8 base/interactiveutil.jl
   6 src/flisp/Makefile
   6 README.md
   5 test/sparse.jl
   5 test/file.jl
   5 src/sys.c
   5 src/support/Makefile
   5 src/debuginfo.cpp
   5 LICENSE.md
   5 contrib/windows/msys_build.sh
   5 appveyor.yml
   5 .travis.yml
```

subject of this talk

Julia's dependencies

Open source and the shoulders of giants

"If I have seen further it is by standing on the shoulders of giants." – Isaac Newton

"We all build our ideas on the best ideas we can find. Now imagine if there were no more good ideas we were allowed to use." – Bob Young, co-founder of Red Hat

- `git clone git://github.com/JuliaLang/julia.git`
`cd julia && make` takes a while – the first time
- `deps/Makefile` downloads, configures, and compiles each dependency – a few of them are much bigger than Julia is
- Three groups of dependencies:
 - 1 Linked into `libjulia`, necessary to run the Julia JIT compiler
 - 2 Used in Julia's standard library via `ccall()` or process spawning
 - 3 Tools used only at build time to compile Julia (or other deps)

What does all this code do?

Group 1, required for libjulia

- LLVM
 - ▶ JIT compiler
- libuv
 - ▶ Cross-platform input/output
- libunwind
 - ▶ Handling backtraces
- utf8proc
 - ▶ Unicode processing



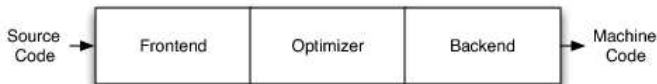
- What: formerly “Low Level Virtual Machine,” today general purpose compiler infrastructure
- Who: many contributors from Apple, Google, Intel, Mozilla, Julia, etc. Used by Clang, Rust, Swift, Emscripten, WebKit (Safari)
- When: originally Chris Lattner’s Masters and Ph.D theses, circa 2003
- License: University of Illinois / NCSA (permissive, BSD-style)
- Written in: C++
- Use in Julia: just in time compiler

A quick introduction to classical compiler design - 1

From “The Architecture of Open Source Applications,”

<http://www.aosabook.org/en/llvm.html>

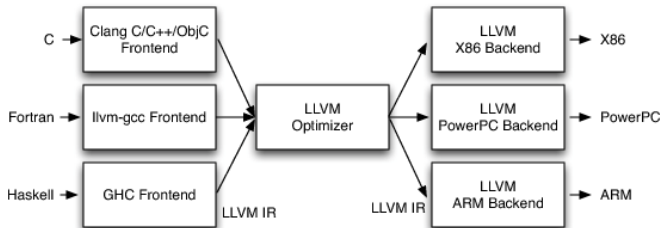
- Basic 3 phase compiler



- One input language, one target architecture

A quick introduction to classical compiler design - 2

- Modular compiler architecture



- Reuse core components across multiple languages and architectures
- LLVM intermediate representation (IR)
 - ▶ Sort of like “cross-platform assembly”
 - ▶ Try out `@code_llvm` in Julia

libuv, <https://github.com/libuv/libuv>



- What: “Lib Unicorn Velociraptor”?
pun on `-luv` to link to the library?
cross platform asynchronous I/O
- Who: originally written to port Node.js
to Windows, also used by Luvit, Pyuv,
formerly Rust
- When: 2011
- License: MIT
- Written in: C
- Use in Julia: input/output, file system,
process spawning, networking, terminal
handling, without requiring POSIX

Julia uses a minor fork of libuv to support process piping syntax
`run(`ls` |> `sort`)`, need work and pull requests to use upstream

libunwind, <http://nongnu.org/libunwind>

- What: Support library for backtraces and profiling on Linux
- Who: David Mosberger, Arun Sharma, other contributors
- When: 2002
- License: MIT
- Written in: C

Different library from Apple, written in C++ and assembly, used on OSX
<https://github.com/JuliaLang/libosxunwind>

utf8proc, <http://public-software-group.org/utf8proc>

- What: Unicode processing, normalization, case folding
- Who: Public Software Group
- When: 2006
- License: MIT
- Written in: C

Unicode is complicated!

- Variable width encodings
- Combining characters
- Different normalization forms
- Similar looking but distinct codepoints

Forked and updated to latest Unicode version by Steven G. Johnson and Jiahao Chen for Julia 0.4-dev, <https://github.com/JuliaLang/libmojibake>

Group 2, used in Julia standard library

- OpenBLAS: Dense linear algebra
- SuiteSparse: Sparse matrix factorizations
- ARPACK: Sparse eigenvalue problems
- FFTW: Fourier transforms
- GMP: Arbitrary precision integers (`BigInt`)
- MPFR: Arbitrary precision floating point (`BigFloat`)
- PCRE: Regular expressions
- OpenLibm: Math functions normally found in `libm`
- OpenSpecFun: Additional special functions
- dSFMT: Random number generation
- Rmath: Distributions and statistical functions
- double-conversion: Floating point to string conversion (Julia $\leq 0.3.x$)
- Git (in Julia $\leq 0.3.x$) or libgit2 (0.4 WIP): Package manager

OpenBLAS, <http://openblas.net>

L	A	P	A	C	K
L	-A	P	-A	C	-K
L	A	P	A	-C	-K
L	-A	P	-A	-C	K
L	A	-P	-A	C	K
L	-A	-P	A	C	-K

- What: Open source high performance implementation of the Basic Linear Algebra Subprograms (BLAS) and Linear Algebra Package (LAPACK)
- Who: originally Kazushige Goto, now maintained by Xianyi Zhang, Werner Saar, and others
- When: GotoBLAS 2002, OpenBLAS 2011
- License: BSD
- Written in: Fortran, assembly, C
- Use in Julia: dense linear algebra

Takes a long time to compile when `OPENBLAS_DYNAMIC_ARCH=1`, builds optimized kernel implementations for many recent CPU families (Nehalem, Sandy Bridge, Haswell, etc)

BLAS in one slide

Level 1 BLAS

	dim	scalar	vector	vector	scalars	5-element array	
SUBROUTINE xROTG (A, B, C, S)		
SUBROUTINE xROTNG(D1, D2, A, B,	PARAM)	
SUBROUTINE xROT (N,			X, INCX, Y, INCY,		C, S)		
SUBROUTINE xROTH (N,			X, INCX, Y, INCY,			PARAM)	
SUBROUTINE xSWAP (N,			X, INCX, Y, INCY)				
SUBROUTINE xSCAL (N,		ALPHA,	X, INCX)				
SUBROUTINE xCOPY (N,			X, INCX, Y, INCY)				
SUBROUTINE xALPY (N,		ALPHA,	X, INCX, Y, INCY)				
FUNCTION xDOT (N,			X, INCX, Y, INCY)				
FUNCTION xDOTU (N,			X, INCX, Y, INCY)				
FUNCTION xDOTC (N,			X, INCX, Y, INCY)				
FUNCTION xxDOT (N,			X, INCX, Y, INCY)				
FUNCTION xNRM2 (N,			X, INCX)				
FUNCTION xASUM (N,			X, INCX)				
FUNCTION IZAMAX (N,			X, INCX)				

Generate plane rotation
 Generate modified plane rotation
 Apply plane rotation
 Apply modified plane rotation
 $x \leftrightarrow y$
 $x \leftarrow \alpha x$
 $y \leftarrow x$
 $y \leftarrow \alpha x + y$
 $\text{dot} \leftarrow x^T y$
 $\text{dot} \leftarrow x^H y$
 $\text{dot} \leftarrow x^T y$
 $\text{dot} \leftarrow x^H y$
 $\text{dot} \leftarrow \alpha + x^T y$
 $\text{nrm2} \leftarrow \|x\|_2$
 $\text{asum} \leftarrow |re(x)|_1 + |im(x)|_1$
 $\text{amax} \leftarrow 1^{st} k \ni |re(x_k)| + |im(x_k)|$
 $\quad = \max(|re(x_i)| + |im(x_i)|)$

prefixes
 S, D
 S, D
 S, D
 S, D
 S, D, C, Z
 S, D, C, Z, CS, ZD
 S, D, C, Z
 S, D, C, Z
 S, D, DS
 C, Z
 C, Z
 SDS
 S, D, SC, DZ
 S, D, SC, DZ
 S, D, C, Z

Level 2 BLAS

options	dim	b-width	scalar	matrix	vector	scalar	vector
xGENV (M, N,	ALPHA, A, LDA,	X, INCX,	BETA, Y, INCY)	
xHEMV (UPLO,			M, N, KL, KU,	ALPHA, A, LDA,	X, INCX,	BETA, Y, INCY)	
xHEMV (UPLO,			M,	ALPHA, A, LDA,	X, INCX,	BETA, Y, INCY)	
xHPMV (UPLO,			M, K,	ALPHA, A, LDA,	X, INCX,	BETA, Y, INCY)	
xSYMV (UPLO,			M,	ALPHA, A, LDA,	X, INCX,	BETA, Y, INCY)	
xSBMV (UPLO,			M, K,	ALPHA, A, LDA,	X, INCX,	BETA, Y, INCY)	
xSPMV (UPLO,			M,	ALPHA, AP,	X, INCX,	BETA, Y, INCY)	
xTRMV (UPLO, TRANS, DIAG,			M,	A, LDA,	X, INCX)		
xTRMV (UPLO, TRANS, DIAG,			M, K,	A, LDA,	X, INCX)		
xTPMV (UPLO, TRANS, DIAG,			M,	AP,	X, INCX)		
xTRSV (UPLO, TRANS, DIAG,			M,	A, LDA,	X, INCX)		
xTRSV (UPLO, TRANS, DIAG,			M, K,	A, LDA,	X, INCX)		
xTSV (UPLO, TRANS, DIAG,			M,	AP,	X, INCX)		
options	dim	scalar	vector	vector	matrix		
xGER (M, N,	ALPHA, X, INCX,	Y, INCY,	A, LDA)		
xGERU (M, N,	ALPHA, X, INCX,	Y, INCY,	A, LDA)		
xGERC (M, N,	ALPHA, X, INCX,	Y, INCY,	A, LDA)		
xHER (UPLO,		M,	ALPHA, X, INCX,	A, LDA)			
xHPR (UPLO,		M,	ALPHA, X, INCX,	AP)			
xHER2 (UPLO,		M,	ALPHA, X, INCX,	Y, INCY,	A, LDA)		
xHPR2 (UPLO,		M,	ALPHA, X, INCX,	Y, INCY,	AP)		
xSYR (UPLO,		M,	ALPHA, X, INCX,	A, LDA)			
xSPR (UPLO,		M,	ALPHA, X, INCX,	AP)			
xSYR2 (UPLO,		M,	ALPHA, X, INCX,	Y, INCY,	A, LDA)		
xSPR2 (UPLO,		M,	ALPHA, X, INCX,	Y, INCY,	AP)		

$y \leftarrow \alpha Ax + \beta y, y \leftarrow \alpha A^T x + \beta y, y \leftarrow \alpha A^H x + \beta y, A - m \times n$
 $y \leftarrow \alpha Ax + \beta y, y \leftarrow \alpha A^T x + \beta y, y \leftarrow \alpha A^H x + \beta y, A - m \times n$
 $y \leftarrow \alpha Ax + \beta y$
 $y \leftarrow \alpha Ax + \beta y$
 $y \leftarrow \alpha Ax + \beta y$
 $y \leftarrow \alpha Ax + \beta y$
 $y \leftarrow \alpha Ax + \beta y$
 $y \leftarrow \alpha Ax + \beta y$
 $x \leftarrow Ax, x \leftarrow A^T x, x \leftarrow A^H x$
 $x \leftarrow Ax, x \leftarrow A^T x, x \leftarrow A^H x$
 $x \leftarrow Ax, x \leftarrow A^T x, x \leftarrow A^H x$
 $x \leftarrow A^{-1} x, x \leftarrow A^{-T} x, x \leftarrow A^{-H} x$
 $x \leftarrow A^{-1} x, x \leftarrow A^{-T} x, x \leftarrow A^{-H} x$
 $x \leftarrow A^{-1} x, x \leftarrow A^{-T} x, x \leftarrow A^{-H} x$

S, D, C, Z
 S, D, C, Z
 C, Z
 C, Z
 C, Z
 S, D
 S, D
 S, D
 S, D, C, Z
 S, D, C, Z
 S, D, C, Z
 S, D, C, Z
 S, D, C, Z
 S, D, C, Z
 S, D, C, Z

$A \leftarrow \alpha xy^T + A, A - m \times n$
 $A \leftarrow \alpha xy^T + A, A - m \times n$
 $A \leftarrow \alpha xy^H + A, A - m \times n$
 $A \leftarrow \alpha xy^H + A$
 $A \leftarrow \alpha xy^H + A$
 $A \leftarrow \alpha xy^H + y(\alpha x)^H + A$
 $A \leftarrow \alpha xy^H + y(\alpha x)^H + A$
 $A \leftarrow \alpha x x^T + A$
 $A \leftarrow \alpha x x^T + A$
 $A \leftarrow \alpha x x^T + A$
 $A \leftarrow \alpha xy^T + \alpha y x^T + A$
 $A \leftarrow \alpha xy^T + \alpha y x^T + A$

S, D
 C, Z
 C, Z
 C, Z
 C, Z
 C, Z
 C, Z
 S, D
 S, D
 S, D
 S, D

Level 3 BLAS

options	dim	scalar	matrix	matrix	scalar	matrix
xGEMM (M, N, K,	ALPHA, A, LDA,	B, LDB,	BETA, C, LDC)	
xSYMM (SIDE, UPLO,		M, N,	ALPHA, A, LDA,	B, LDB,	BETA, C, LDC)	
xHEMM (SIDE, UPLO,		M, N,	ALPHA, A, LDA,	B, LDB,	BETA, C, LDC)	
xSYRK (UPLO, TRANS,		M, K,	ALPHA, A, LDA,	BETA, C, LDC)		
xHERK (UPLO, TRANS,		M, K,	ALPHA, A, LDA,	BETA, C, LDC)		
xSYRK (UPLO, TRANS,		M, K,	ALPHA, A, LDA,	B, LDB,	BETA, C, LDC)	
xHERK (UPLO, TRANS,		M, K,	ALPHA, A, LDA,	B, LDB,	BETA, C, LDC)	
xTRMM (SIDE, UPLO, TRANS,		DIAG, M, N,	ALPHA, A, LDA,	B, LDB)		
xTRSM (SIDE, UPLO, TRANS,		DIAG, M, N,	ALPHA, A, LDA,	B, LDB)		

$C \leftarrow \alpha op(A) op(B) + \beta C, op(X) = X, X^T, X^H, C - m \times n$
 $C \leftarrow \alpha AB + \beta C, C \leftarrow \alpha BA + \beta C, C - m \times n, A = A^T$
 $C \leftarrow \alpha AB + \beta C, C \leftarrow \alpha BA + \beta C, C - m \times n, A = A^H$
 $C \leftarrow \alpha A A^T + \beta C, C \leftarrow \alpha A^H A + \beta C, C - n \times n$
 $C \leftarrow \alpha A A^H + \beta C, C \leftarrow \alpha A^T A + \beta C, C - n \times n$
 $C \leftarrow \alpha A B^T + \beta C, C \leftarrow \alpha A^T B + \beta B^T A + \beta C, C - n \times n$
 $C \leftarrow \alpha A B^H + \beta B A^H + \beta C, C \leftarrow \alpha A^H B + \beta B^H A + \beta C, C - n \times n$
 $B \leftarrow \alpha op(A) B + \alpha B op(A), op(A) = A, A^T, A^H, B - m \times n$
 $B \leftarrow \alpha op(A^{-1}) B, B \leftarrow \alpha B op(A^{-1}), op(A) = A, A^T, A^H, B - m \times n$

S, D, C, Z
 S, D, C, Z
 C, Z
 S, D, C, Z
 C, Z
 S, D, C, Z
 S, D, C, Z
 S, D, C, Z
 S, D, C, Z

Levels of BLAS and LAPACK

- Level 1 BLAS
 - ▶ Vector operations: scale, add, copy, dot product, find largest element
 - ▶ $O(n)$ operations on $O(n)$ data
- Level 2 BLAS
 - ▶ Matrix-vector operations: $A * x$, triangular solve, rank 1 or 2 updates
 - ▶ $O(n^2)$ operations on $O(n^2)$ data
- Level 3 BLAS
 - ▶ Matrix-matrix operations: $A * B$, triangular solve for multiple right hand sides, rank k updates
 - ▶ $O(n^3)$ operations on $O(n^2)$ data
 - ▶ Where the important cache optimizations happen
- LAPACK (API does not fit on one slide)
 - ▶ Higher level factorizations, eigenvalue and singular value decompositions designed to use efficient BLAS-3 operations
- Reference Fortran implementations from Netlib are slow
 - ▶ Better to use an optimized (SIMD, multithreaded) implementation like OpenBLAS or Intel MKL

SuiteSparse, <http://suitsparse.com>



- What: Sparse linear algebra – LU, Cholesky, and QR factorizations
- Who: Tim Davis, also used by Matlab, Mathematica, Google Ceres
- When: 2006 or earlier?
- License: LGPL
- Written in: C, C++

What makes a matrix sparse?

- Enough elements are zero to be worth taking advantage of
- Nonzero structure and permutations very important
- Graph theory for structure, linear algebra for numerics

ARPACK, <https://github.com/opencollab/arpack-ng>

- What: ARnoldi PACKAge for eigenvalue problems $Ax = \lambda x$ with sparse A , iterative algorithms to find a small set of λ values
- Who: Rich Lehoucq, Kristi Maschhoff, et al, also used by SciPy, Matlab, Octave
- When: 1996
- License: BSD
- Written in: Fortran 77

Work in progress to replace this with pure Julia code, see <https://github.com/JuliaLang/IterativeSolvers.jl/pull/31>



- What: Fastest Fourier Transform in the West
- Who: Matteo Frigo and Steven G. Johnson
- When: 1997
- License: GPL
- Written in: C, code generator in OCaml

Work in progress to write a pure Julia FFT and move FFTW to an optional package, see <https://github.com/JuliaLang/julia/pull/6193>



- What: GNU Multiple Precision library for arbitrary precision integer arithmetic (BigInt in Julia)
- Who: Torbjörn Granlund, GNU Project, many contributors
- When: 1991
- License: LGPL
- Written in: C, assembly



- What: GNU Multiple Precision Floating-point Reliable library for arbitrary precision floating-point arithmetic (BigFloat in Julia)
- Who: INRIA, GNU Project, many contributors
- When: 2000
- License: LGPL
- Written in: C

Some people, when confronted with a problem, think “I know, I’ll use regular expressions.” Now they have two problems. – Jamie Zawinski

- What: Perl Compatible Regular Expressions (r"..." in Julia)
- Who: Philip Hazel, Zoltan Herczeg
- When: 1997
- License: BSD
- Written in: C

OpenLibm, <http://openlibm.org>

- What: “high quality, portable, standalone C mathematical library”
- Who: Viral Shah and other JuliaLang contributors, code originally from FreeBSD msun, OpenBSD libm, and Freely Distributable FDLIBM (<http://www.netlib.org/fdlibm>)
- When: 2011, original code from 1992 or earlier
- License: BSD
- Written in: C, assembly
- Why: Performance and accuracy of system `libm` for `trig`, `sqrt`, `exp`, `log`, etc varies between platforms, more reliable to build our own

OpenSpecFun, <https://github.com/JuliaLang/openspecfun>

- What: Collection of special functions – Bessel and Airy functions from AMOS library, complex error functions from Faddeeva
- Who: Donald Amos from <http://netlib.org/amos>, Faddeeva by Steven G. Johnson
http://ab-initio.mit.edu/wiki/index.php/Faddeeva_Package
- When: Amos code from 1985 or earlier, Faddeeva 2012, OpenSpecFun 2013
- License: MIT, Public Domain
- Written in: Fortran 77, C++, C
- Why: Pieces that were split out from OpenLibm because they are not included in system libm libraries

dSFMT, <http://www.math.sci.hiroshima-u.ac.jp/~m-mat/MT/SFMT>

- What: Double precision SIMD-oriented Fast Mersenne Twister random number generator (RNG)
- Who: Matsuo Saito and Makoto Matsumoto
- When: 2007
- License: BSD
- Written in: C

Julia's RNG code is undergoing heavy development on 0.4-dev right now, good chance dSFMT might be replaced by some alternate RNG
see <https://github.com/JuliaLang/julia/issues/8786>

Rmath, <https://github.com/JuliaLang/Rmath>

- What: Distributions and statistical functions from R, forked and patched to use same dSFMT RNG as Julia
- Who: R contributors, fork by Viral Shah
- When: Possibly as old as R, 1997? JuliaLang fork from 2013
- License: GPL
- Written in: C

Not used by base Julia, only Distributions.jl and HypothesisTests.jl!
see <https://github.com/JuliaStats/Distributions.jl/pull/138>
for work to remove the dependency